


Modern Heterogeneous Computers:

(4) Memory Components II

Hung-Wei Tseng

Recap: the roofline after using hardware accelerators



Memory technologies we have today

Volatile Memory

Non-Volatile Memory

100ps

SRAM

ns

DRAM

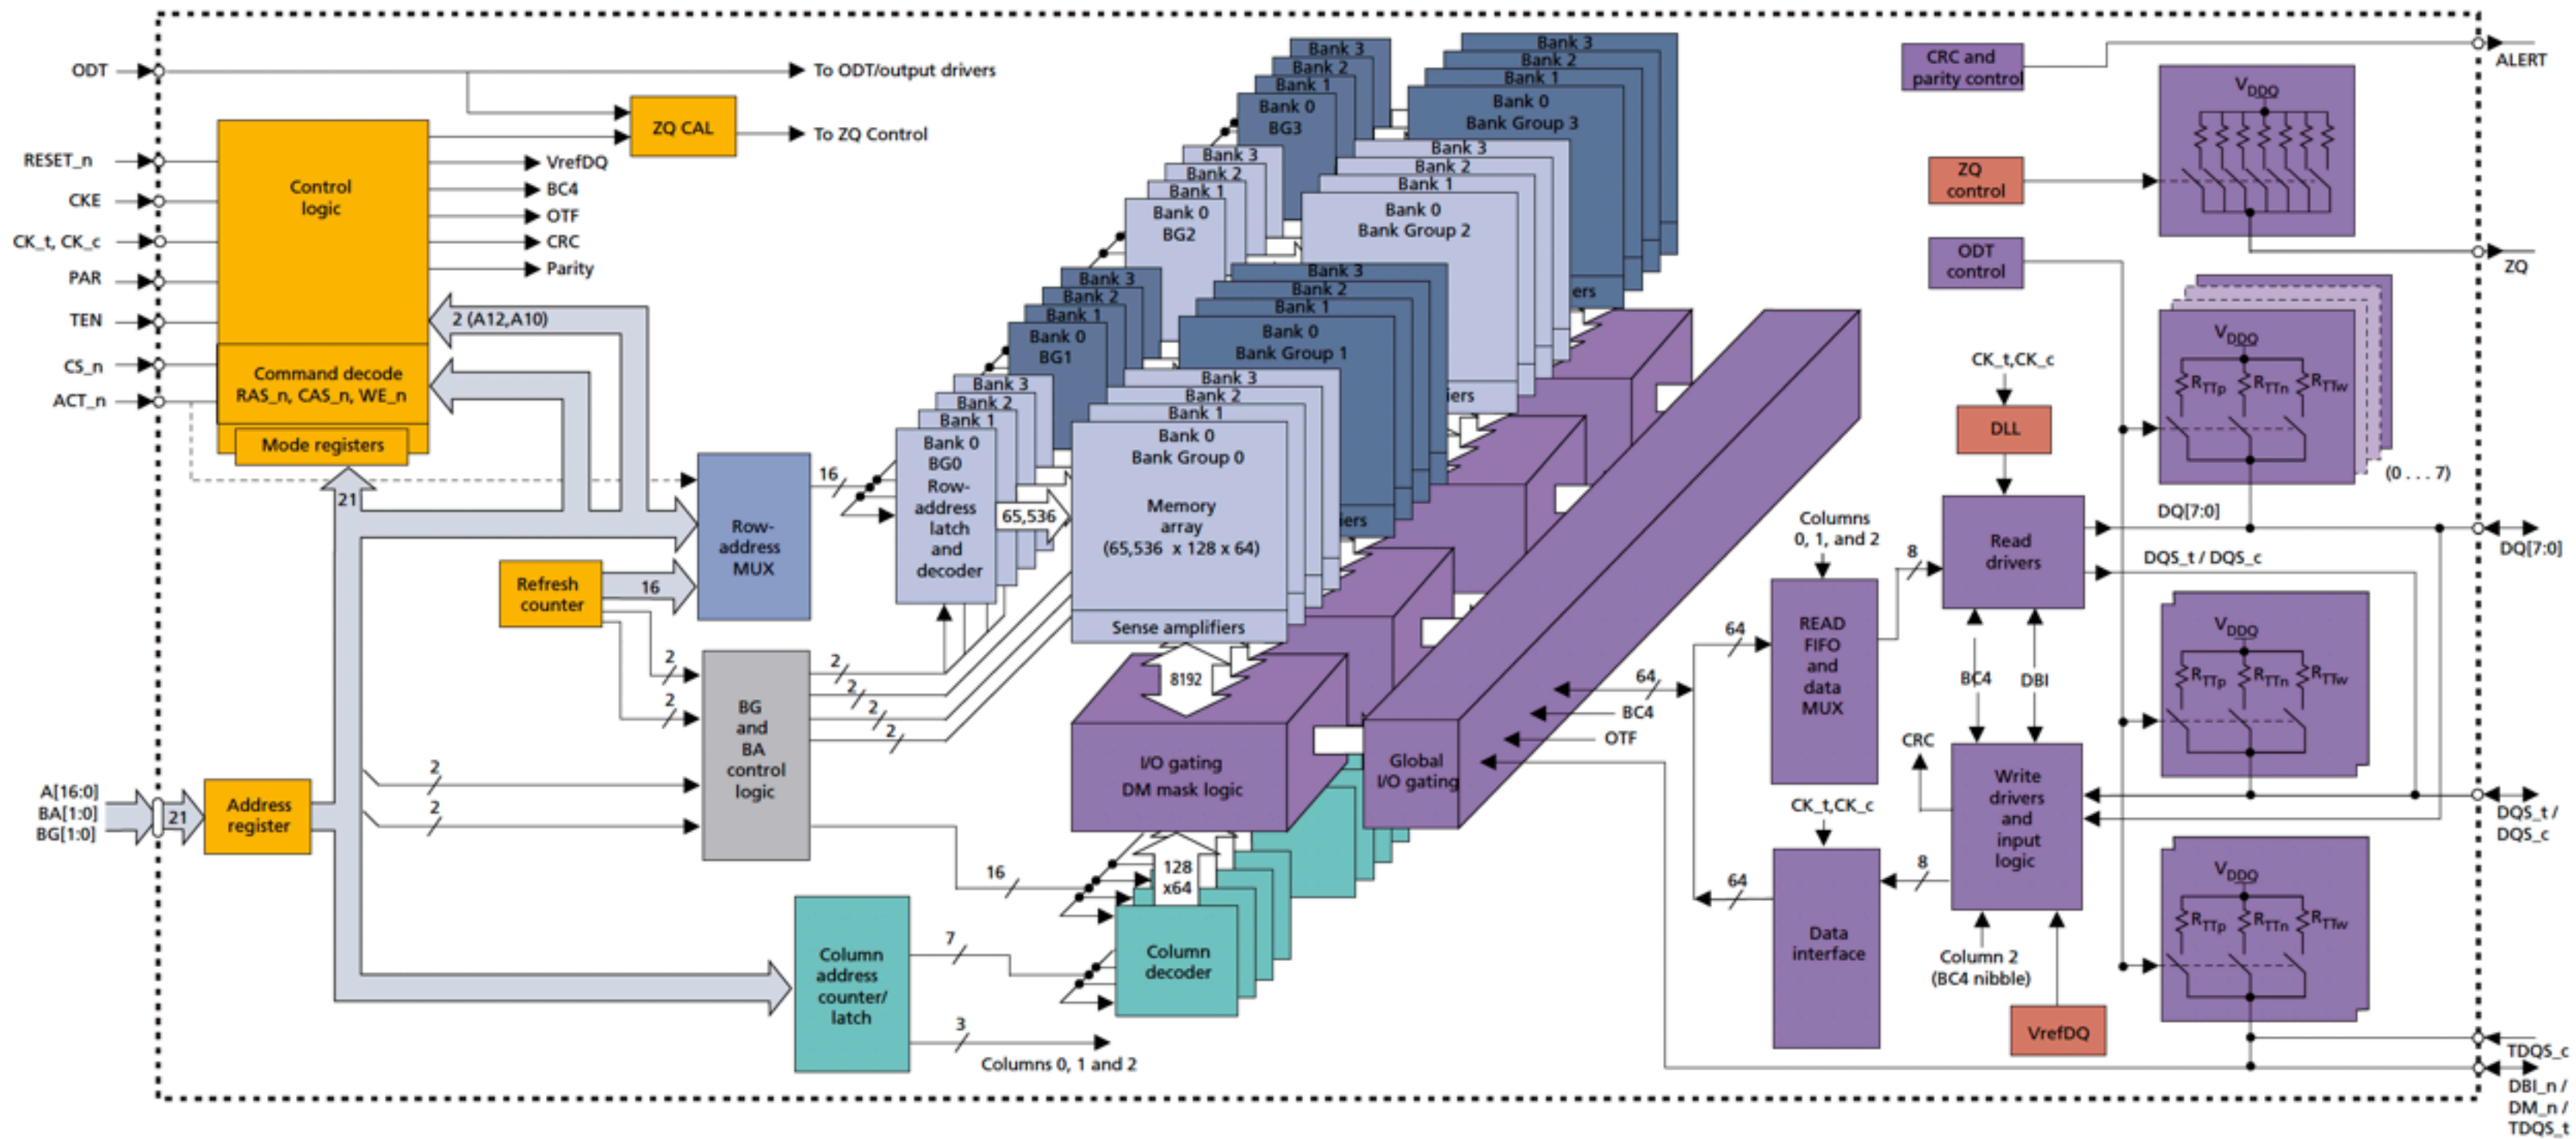
RRAM

us


**PCM 3DXPoint
Flash memory**

ms

Hard Disk Drives



Recap: DRAM Accesses



The latency of pre-charge, row/column accesses is fully covered!

2. Key Features

[Table 2] 8Gb DDR4 C-die Speed bins

Speed	DDR4-1600	DDR4-1866	DDR4-2133	DDR4-2400	DDR4-2666	Unit
	11-11-11	13-13-13	15-15-15	17-17-17	19-19-19	
tCK(min)	1.25	1.071	0.937	0.833	0.75	ns
CAS Latency	11	13	15	17	19	nCK
tRCD(min)	13.75	13.92	14.06	14.16	14.25	ns
tRP(min)	13.75	13.92	14.06	14.16	14.25	ns
tRAS(min)	35	34	33	32	32	ns
tRC(min)	48.75	47.92	47.06	46.16	46.25	ns

• JEDEC standard 1.2V (1.14V~1.26V)

The 8Gb DDR4 SDRAM C-die is organized as a 64Mbit x 16 I/Os x 8banks

Recap: DRAM Performance

2. Key Features

- Latency per "8-bit" — 0.75 ns (if it's row-buffered)

- Bandwidth per die = $\frac{1}{0.75ns} = 1.33GB/sec$

- 16 chips = $16 \times \frac{1}{0.75ns} = 21.33GB/sec$

[Table 2] 8Gb DDR4 C-die Speed bins

Speed	DDR4-1600	DDR4-1866	DDR4-2133	DDR4-2400	DDR4-2666
	11-11-11	13-13-13	15-15-15	17-17-17	19-19-19
tCK(min)	1.25	1.071	0.937	0.833	0.75
CAS Latency	11	13	15	17	19
tRCD(min)	13.75	13.92	14.06	14.16	14.25
tRP(min)	13.75	13.92	14.06	14.16	14.25
tRAS(min)	35	34	33	32	32
tRC(min)	48.75	47.92	47.06	46.16	46.25

- JEDEC standard 1.2V (1.14V~1.26V)
- V_{DDQ} = 1.2V (1.14V~1.26V)
- V_{PP} = 2.5V (2.375V~2.75V)
- 800 MHz f_{CK} for 1600Mb/sec/pin, 933 MHz f_{CK} for 1866Mb/sec/pin, 1067MHz f_{CK} for 2133Mb/sec/pin, 1200MHz f_{CK} for 2400Mb/sec/pin, 1333MHz f_{CK} for 2666Mb/sec/pin
- 8 Banks (2 Bank Groups)
- Programmable CAS Latency (posted CAS): 10,11,12,13,14,15,16,17,18,19,20
- Programmable CAS Write Latency (CWL) = 9,11 (DDR4-1600), 10,12 (DDR4-1866),11,14 (DDR4-2133),12,16 (DDR4-2400) and 14,18 (DDR4-2666)
- 8-bit pre-fetch
- Burst Length: 8, 4 with tCCD = 4 which does not allow seamless read or write [either On the fly using A12 or MRS]
- Bi-directional Differential Data-Strobe
- Internal (self) calibration: Internal self calibration through ZQ pin (RZQ: 240 ohm ± 1%)
- On Die Termination using ODT pin
- Average Refresh Period 7.8us at lower than T_{CASE} 85°C, 3.9us at 85°C < T_{CASE} ≤ 95 °C
- Connectivity Test Mode (TEN) is Supported


The 8Gb DDR4 SDRAM C-die is organized as a 64Mbit x device. This synchronous device achieves high speed transfer rates of up to 2666Mb/sec/pin (DDR4-2666) for tions.

The chip is designed to comply with the following key DRAM features such as posted CAS, Programmable CWL, Internal On Die Termination using ODT pin and Asynchronous Read

All of the control and address inputs are synchronized with internally supplied differential clocks. Inputs are latched at the differential clocks (CK rising and \overline{CK} falling). All I/Os are synchronous pair of bidirectional strobes (DQS and \overline{DQS}) in a source synchronous mode. The address bus is used to convey row, column, and bank information in a $\overline{RAS}/\overline{CAS}$ multiplexing style. The DDR4 device is powered with a single 1.2V (1.14V~1.26V) power supply, 1.2V(1.14V~1.26V) and 2.5V (2.375V~2.75V) V_{PP}.

The 8Gb DDR4 C-die device is available in 96ball FBGA package.

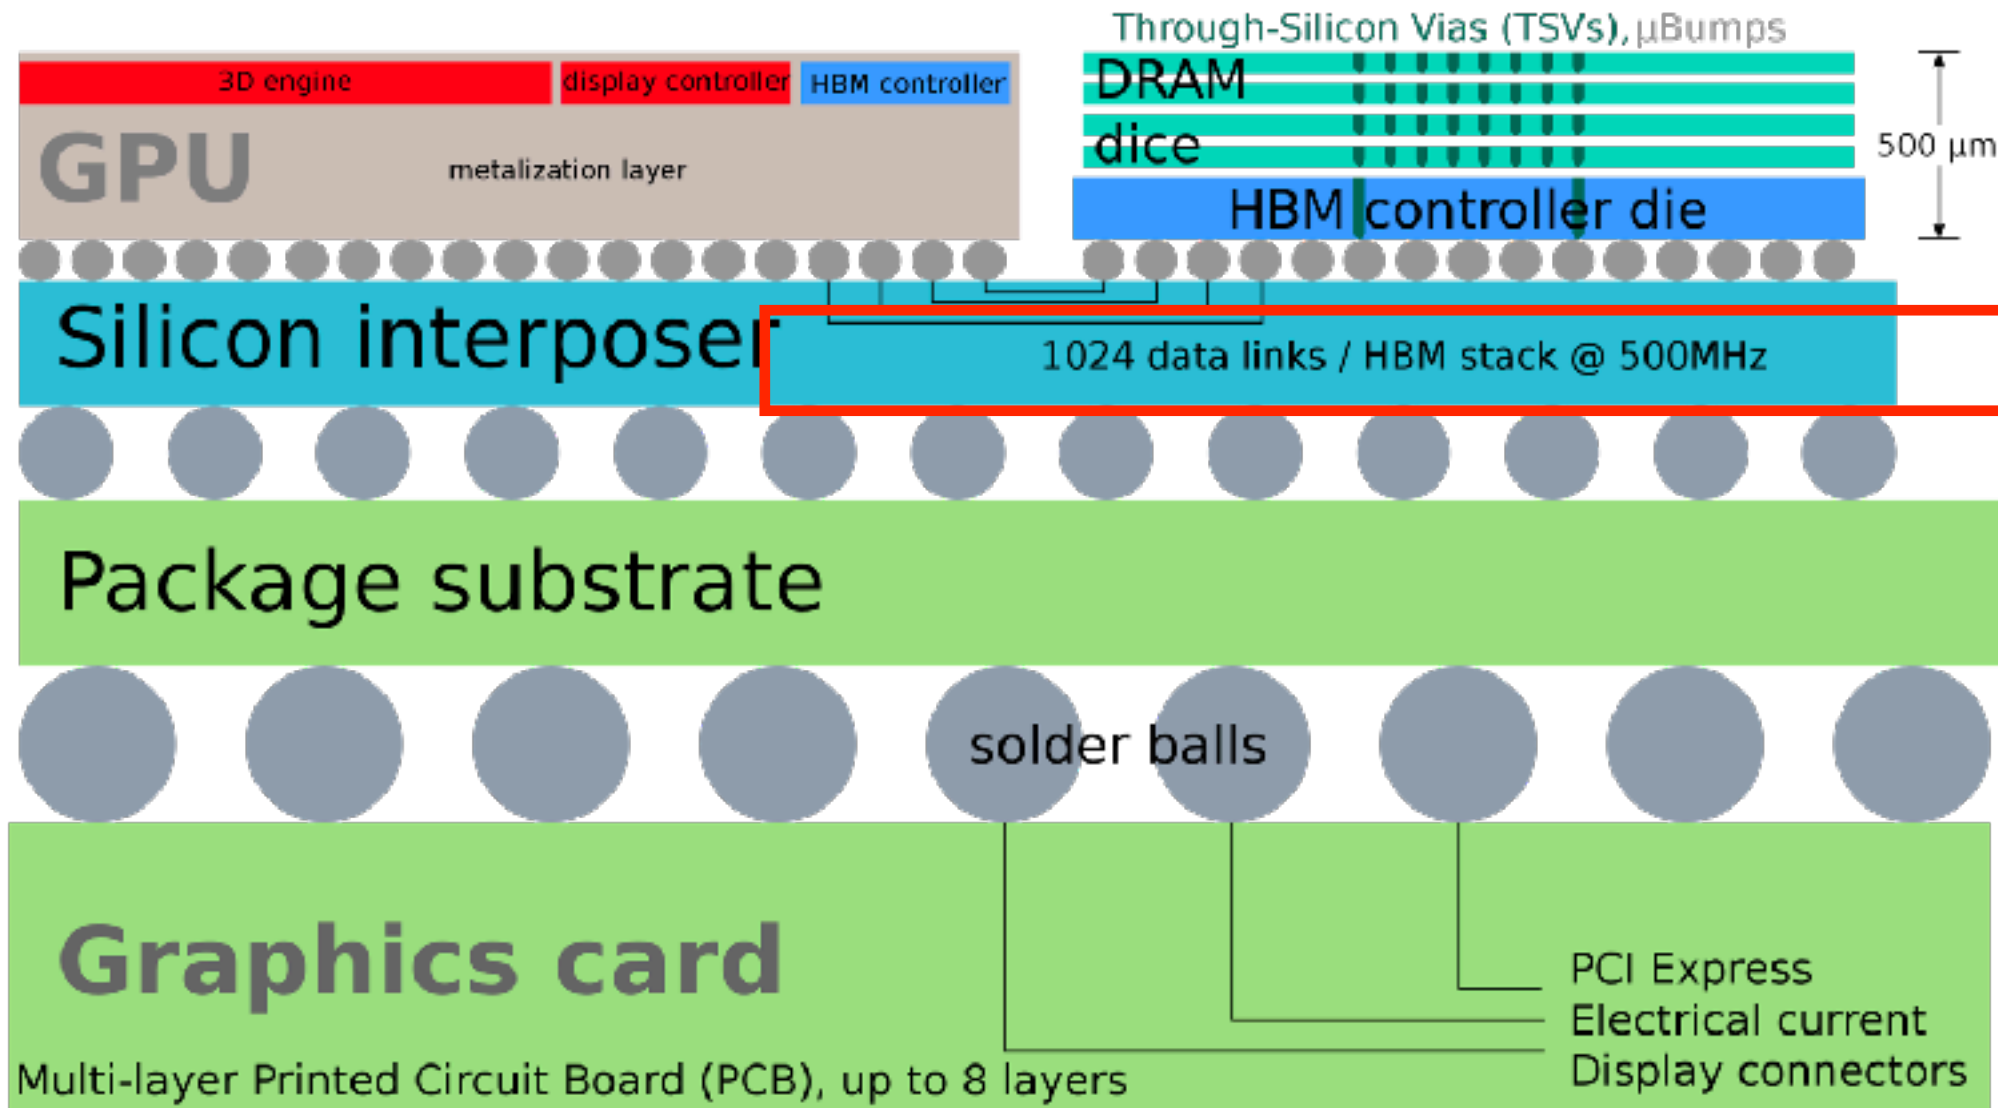
Recap: the roofline after using hardware accelerators



Ideas of increasing bandwidth

- More parallel bits
- More banks
- More channels
- Widen the memory-processor bus

HBM (High Bandwidth Memory)




$$0.5G \times \frac{1024bits}{8bits} = 64GB/sec$$

If you have 4x modules —
 $4 * 64 = 256 GB/sec$

HBM2 increase the clock
rate to 2GHz — 1TB/sec

Recap: the roofline after using hardware accelerators



What's the pros & cons of GDDR v.s. HBM

GDDR v.s. HBM

GDDR v.s. HBM

- HBM's bandwidth is wider
- HBM is more expensive — pin counts are more expensive
- HBM is limited by the per-chip heat dissipation — less powerful

Which one do you prefer?

**A. Single chip, heterogeneous processor
(e.g., NVIDIA/intel) or**

B. standalone accelerators? (e.g., TPU)

Single-chip or standalone ones?

Is system-wide heterogeneous processing a better idea?

- Easier for heat dissipation
 - Each processor can be more powerful
 - gamers/datacenters still prefers discrete GPUs
 - Cloud TPUs are standalone ones
- System size is larger, cost of ownership can be higher
- Data movement going through system interconnects
- Allow the “computation core” of the accelerator to enjoy larger bandwidth, if you do it right.

Case: what's the goal of the following program and what do you expect the performance would change when size varies?

3:00

```
#include <stdio.h>
#include <stdlib.h>
#include <sys/time.h>
#include <time.h> /* for clock_gettime */
#include <malloc.h>
#include <immintrin.h>

void write_memory_avx(void* array, size_t size) {
    __m256i* varray = (__m256i*) array;

    __m256i vals = _mm256_set1_epi32(1);
    size_t i;
    for (i = 0; i < size / sizeof(__m256i); i++) {
        _mm256_store_si256(&varray[i], vals); // This will
        generate the vmovaps instruction.
    }
}

int main(int argc, char **argv)
{
    size_t *array;
    size_t size;
    double total_time;
    struct timespec start, end;
    int i;
    size = atoi(argv[1])/sizeof(size_t);
```

```
    array = (size_t *)memalign(32, sizeof(size_t)*size);
    clock_gettime(CLOCK_MONOTONIC, &start); /* mark start
time */
    write_memory_avx(array, size);

    clock_gettime(CLOCK_MONOTONIC, &end); /* mark the end
time */
    total_time = ((end.tv_sec * 1000000000.0 + end.tv_nsec)
- (start.tv_sec * 1000000000.0 + start.tv_nsec));
    fprintf(stderr, "Latency: %.0lf ns, GBps:
%.0lf\n", total_time, (double)((double)size*sizeof(size_t)/
(total_time)));
    clock_gettime(CLOCK_MONOTONIC, &start); /* mark start
time */
    for(i = 0 ; i< 10; i++)
        write_memory_avx(array, size);

    clock_gettime(CLOCK_MONOTONIC, &end); /* mark the end
time */
    total_time = ((end.tv_sec * 1000000000.0 + end.tv_nsec)
- (start.tv_sec * 1000000000.0 + start.tv_nsec));
    fprintf(stderr, "Latency (10x average): %.0lf ns, GBps
(10x average): %.0lf\n", total_time/10, (double)
((double)size*10*sizeof(size_t)/(total_time)));
    return 0;
}
```

<https://github.com/hungweitseng/EE277/tree/main/demo/memory>

The program

```
#include <stdio.h>
#include <stdlib.h>
#include <sys/time.h>
#include <time.h> /* for clock_gettime */
#include <malloc.h>
#include <immintrin.h>

void write_memory_avx(void* array, size_t size) {
    __m256i* varray = (__m256i*) array;

    __m256i vals = _mm256_set1_epi32(1);
    size_t i;
    for (i = 0; i < size / sizeof(__m256i); i++) {
        _mm256_store_si256(&varray[i], vals); // This will
        generate the vmovaps instruction.
    }
}

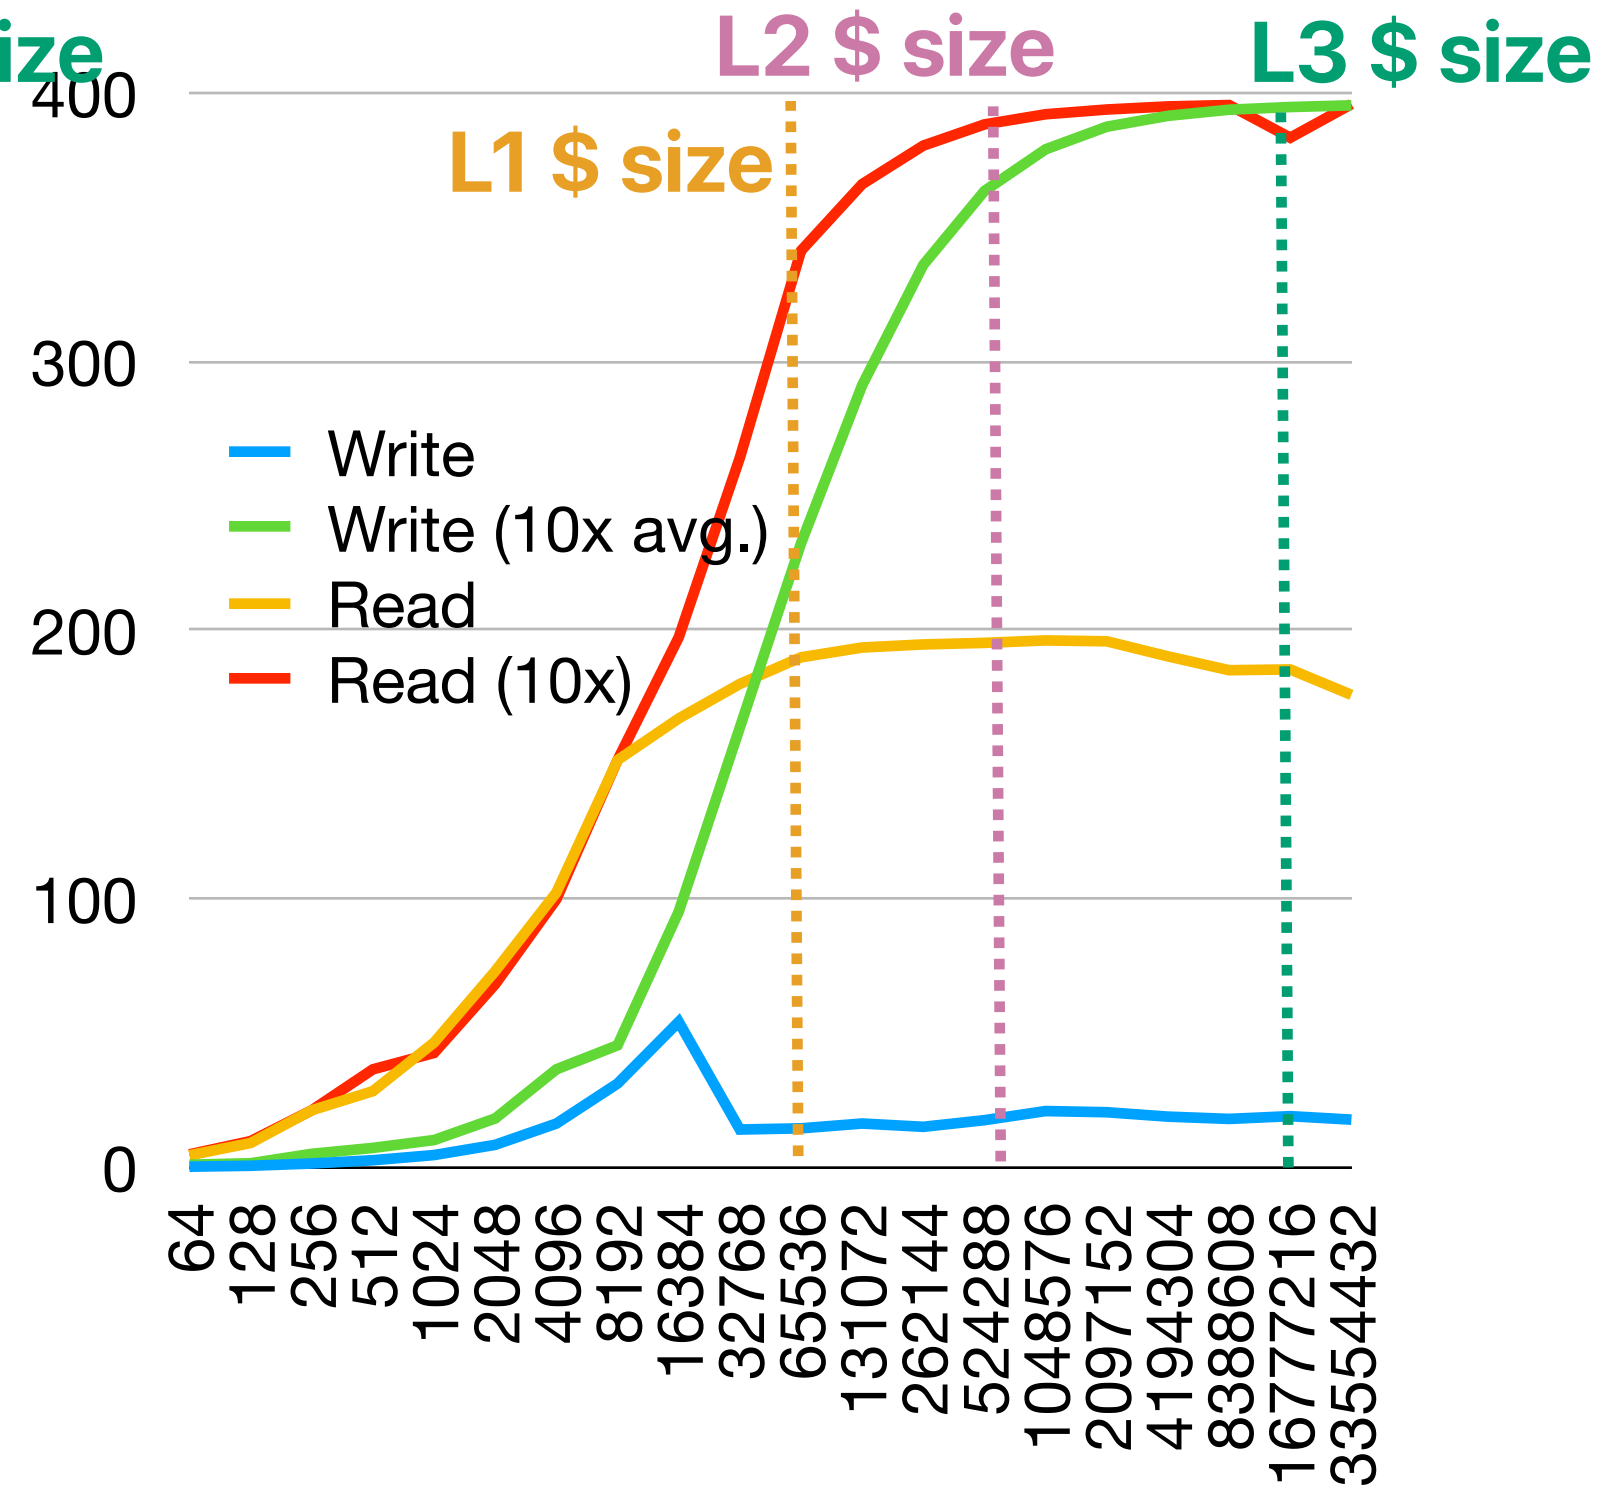
int main(int argc, char **argv)
{
    size_t *array;
    size_t size;
    double total_time;
    struct timespec start, end;
    int i;
    size = atoi(argv[1])/sizeof(size_t);
    array = (size_t *)memalign(32,sizeof(size_t)*size);
```

```
    clock_gettime(CLOCK_MONOTONIC, &start); /* mark start
time */
    write_memory_avx(array, size);

    clock_gettime(CLOCK_MONOTONIC, &end); /* mark the end
time */
    total_time = ((end.tv_sec * 1000000000.0 + end.tv_nsec)
- (start.tv_sec * 1000000000.0 + start.tv_nsec));
    fprintf(stderr,"Latency: %.0lf ns, GBps:
%.0lf\n",total_time, (double)((double)size*sizeof(size_t)/
(total_time)));
    clock_gettime(CLOCK_MONOTONIC, &start); /* mark start
time */
    for(i = 0 ; i< 10; i++)
        write_memory_avx(array, size);

    clock_gettime(CLOCK_MONOTONIC, &end); /* mark the end
time */
    total_time = ((end.tv_sec * 1000000000.0 + end.tv_nsec)
- (start.tv_sec * 1000000000.0 + start.tv_nsec));
    fprintf(stderr,"Latency (10x average): %.0lf ns, GBps
(10x average): %.0lf\n",total_time/10, (double)
((double)size*10*sizeof(size_t)/(total_time)));
    return 0;
}
```

Performance Chart (on AMD Ryzen 5 2600)



3D-die stacking

- Providing huge bandwidth between memory chips and processors (e.g., HBM)
- The renaissance of near-memory processing!
 - Zhu, Qiuling, Berkin Akin, H. Ekin Sumbul, Fazle Sadi, James C. Hoe, Larry Pileggi, and Franz Franchetti. A 3D-stacked logic-in-memory accelerator for application-specific data intensive computing. In 3DIC. 2013.
 - Dongping Zhang, Nuwan Jayasena, Alexander Lyashevsky, Joseph L. Greathouse, Lifan Xu, and Michael Ignatowski. TOP-PIM: throughput-oriented programmable processing in memory. HPDC '14. 2014
 - Farmahini-Farahani, Amin, Jung Ho Ahn, Katherine Morrow, and Nam Sung Kim. NDA: Near-DRAM acceleration architecture leveraging commodity DRAM devices and standard memory modules. In HPCA. 2015.
 - Junwhan Ahn, Sungpack Hong, Sungjoo Yoo, Onur Mutlu, and Kiyoun Choi. A scalable processing-in-memory accelerator for parallel graph processing. ISCA '15. 2015.
 - Youngeun Kwon, Yunjae Lee, and Minsoo Rhu. TensorDIMM: A Practical Near-Memory Processing Architecture for Embeddings and Tensor Operations in Deep Learning. In MICRO '52. 2019

Electrical Computer Science Engineering

277

つくづく

